

```
# For data loading and handling
import pandas as pd
import numpy as np

# For text preprocessing
import re #Removes unwanted patterns (punctuation, numbers).
import string #Provides punctuation symbols.
import nltk #Used for text preprocessing.
from nltk.corpus import stopwords #Removes common words like is, the, and.
from nltk.stem import PorterStemmer #Reduces words to root form

# For feature extraction
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

# For train-test split
from sklearn.model_selection import train_test_split

# For model building
from sklearn.naive_bayes import MultinomialNB

# For evaluation
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Download stopwords (only once) - ensures necessary data for NLP is available
nltk.download('stopwords')
```

[nltk\_data] Downloading package stopwords to /root/nltk\_data...
[nltk\_data] Package stopwords is already up-to-date!
True

```
# Load dataset from the specified CSV file
df = pd.read_csv('/content/news.csv')

# Display the first 5 rows of the DataFrame to inspect the data structure
print("First 5 samples:")
print(df.head())

# Print the shape of the dataset (number of rows, number of columns)
print("\nDataset Shape:", df.shape)

# Display the column names to understand available features
print("\nColumns:", df.columns)

# Assuming:
# text column = 'text' (This is the content of the news article)
# label column = 'label' (This indicates if the news is 'FAKE' or 'REAL')
# (Change column names if different based on your dataset)

# Print the distribution of classes in the 'label' column to check for imbalance
print("\nClass Distribution:")
print(df['label'].value_counts())
```

```
First 5 samples:
      Unnamed: 0          title \
0        8476    You Can Smell Hillary's Fear
1      10294  Watch The Exact Moment Paul Ryan Committed Pol...
2       3608      Kerry to go to Paris in gesture of sympathy
3      10142  Bernie supporters on Twitter erupt in anger ag...
4        875  The Battle of New York: Why This Primary Matters
```

```
           text label
0  Daniel Greenfield, a Shillman Journalism Fello...  FAKE
1  Google Pinterest Digg LinkedIn Reddit Stumbleu...  FAKE
2  U.S. Secretary of State John F. Kerry said Mon...  REAL
3  – Kaydee King (@KaydeeKing) November 9, 2016 T...  FAKE
4  It's primary day in New York and front-runners...  REAL
```

Dataset Shape: (6335, 4)

Columns: Index(['Unnamed: 0', 'title', 'text', 'label'], dtype='object')

Class Distribution:

label	
REAL	3171
FAKE	3164

Name: count, dtype: int64

```
# =====
# STEP 4: Text Preprocessing
# =====
```

```
# Initialize stop words for English language
stop_words = set(stopwords.words('english'))
# Initialize Porter Stemmer for reducing words to their root form
stemmer = PorterStemmer()

# Define a function to preprocess text
def preprocess_text(text):
    # Convert text to lowercase
    text = text.lower()

    # Remove punctuation from the text using a regular expression
    text = re.sub(f"[{string.punctuation}]", "", text)

    # Remove numbers from the text
    text = re.sub(r'\d+', '', text)

    # Tokenize the text into individual words
    words = text.split()

    # Remove stopwords and apply stemming to each word
    words = [stemmer.stem(word) for word in words if word not in stop_words]

    # Join the processed words back into a single string
    return " ".join(words)

# Apply the preprocessing function to the 'text' column and store in a new 'clean_text' column
df['clean_text'] = df['text'].apply(preprocess_text)

# Display a sample of the original and cleaned text to verify preprocessing
print("\nSample Cleaned Text:")
print(df[['text', 'clean_text']].head())
```

Sample Cleaned Text:

	text	clean_text
0	Daniel Greenfield, a Shillman Journalism Fello...	daniel greenfield shillman journal fellow free...
1	Google Pinterest Digg Linkedin Reddit Stumbleu...	googl pinterest digg linkedin reddit stumbleup...
2	U.S. Secretary of State John F. Kerry said Mon...	us secretari state john f kerri said monday st...
3	- Kaydee King (@KaydeeKing) November 9, 2016 T...	- kayde king kaydeek novemb lesson tonight dem...
4	It's primary day in New York and front-runners...	primari day new york frontrunn hillari clinton...

```
# Initialize TF-IDF Vectorizer with a maximum of 5000 features
vectorizer = TfidfVectorizer(max_features=5000)

# Fit the vectorizer on the 'clean_text' and transform it into a feature matrix X
X = vectorizer.fit_transform(df['clean_text'])

# Assign the 'label' column to y, which will be our target variable
y = df['label']

# Print the shape of the feature matrix (number of samples, number of features)
print("Feature Matrix Shape:", X.shape)

# Display the first 20 feature names extracted by the vectorizer
print("\nSample Feature Names:")
print(vectorizer.get_feature_names_out()[:20])
```

Feature Matrix Shape: (6335, 5000)

Sample Feature Names:

['abandon' 'abc' 'abdullah' 'abedin' 'abil' 'abl' 'aboard' 'abolish'
'aborigin' 'abort' 'about' 'abraham' 'abroad' 'absenc' 'absent' 'absente'
'absolut' 'absorb' 'absurd' 'abu']

```
# Split the data into training and testing sets
# X: features, y: target labels
# test_size=0.2: 20% of data for testing, 80% for training
# random_state=42: ensures reproducibility of the split
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=42
)

# Print the shape of the training feature set
print("Training size:", X_train.shape)
# Print the shape of the testing feature set
print("Testing size:", X_test.shape)
```

```
print("Training size: ", X_train.shape)
```

```
Training size: (5068, 5000)
```

```
print("Testing size: ", X_test.shape)
```

```
# Initialize the Multinomial Naive Bayes model
model = MultinomialNB()
```

```
# Train the model using the training features (X_train) and training labels (y_train)
model.fit(X_train, y_train)
```

```
# Confirm that the model training is complete
print("Model trained successfully!")
```

```
# Display the parameters of the trained model
print("\nModel Parameters:")
print(model.get_params())
```

```
Model trained successfully!
```

```
Model Parameters:
```

```
{'alpha': 1.0, 'class_prior': None, 'fit_prior': True, 'force_alpha': True}
```

```
# Make predictions on the test set using the trained model
y_pred = model.predict(X_test)
```

```
# Calculate and print the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

```
# Generate and print the classification report, which includes precision, recall, f1-score
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

```
# Generate and print the confusion matrix to evaluate classification performance
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

```
Accuracy: 0.8887134964483031
```

```
Classification Report:
```

	precision	recall	f1-score	support
FAKE	0.88	0.90	0.89	628
REAL	0.90	0.88	0.89	639
accuracy			0.89	1267
macro avg	0.89	0.89	0.89	1267
weighted avg	0.89	0.89	0.89	1267

```
Confusion Matrix:
```

```
[[566  62]
 [ 79 560]]
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Get the confusion matrix from the last step (calculated in the previous cell)
cm = confusion_matrix(y_test, y_pred)
```

```
# Define class labels for the confusion matrix for better readability
class_labels = ['FAKE', 'REAL']
```

```
# Create a figure with a specific size for the heatmap
plt.figure(figsize=(8, 6))
```

```
# Generate the heatmap using seaborn, with annotations, integer format, and 'Blues' colormap
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_labels, yticklabels=class_labels)
```

```
# Set the title of the plot
plt.title('Confusion Matrix')
```

```
# Set the label for the x-axis (Predicted Label)
plt.xlabel('Predicted Label')
```

```
# Set the label for the y-axis (True Label)
plt.ylabel('True Label')
```

```
# Display the plot
plt.show()
```

